

Principal Component Analysis

Machine Learning 1 — Lecture 6

23th April 2024

Robert Peharz

Institute of Theoretical Computer Science

Graz University of Technology

- So far, we considered **linear models** for **supervised learning**
- That is, we predicted some **target** y from an **input vector** \mathbf{x} :
 - **linear regression** (for regression)
 - **logistic regression** (for classification)
 - both can be made non-linear, by replacing \mathbf{x} with a non-linear transformation $\phi(\mathbf{x}): \mathbb{R}^D \mapsto \mathbb{R}^K$
- Today, we will use a linear model for an **unsupervised learning** technique: **principal component analysis (PCA)**
- Thus, we will have a data set

$$\mathcal{D} = \left\{ \mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)} \right\}$$

of input vectors, but **no target**

- Goal of unsupervised learning: find “interesting structure” in the data

PCA is a technique for **dimensionality reduction**.

Generally, dimensionality reduction means that we replace our data

$$\mathcal{D} = \left\{ \mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)} \right\}$$

containing D -dimensional vectors with a new dataset

$$\mathcal{D}' = \left\{ \mathbf{z}^{(1)}, \mathbf{z}^{(2)}, \dots, \mathbf{z}^{(N)} \right\}$$

of K -dimensional vectors, $K \ll D$, where \mathcal{D}' captures “certain characteristics” of \mathcal{D} .

Replace $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ with $\mathcal{D}' = \{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(N)}\}$.


Potential target characteristics (qualitative):

- if $\mathbf{x}^{(i)}$ and $\mathbf{x}^{(j)}$ are close, then also $\mathbf{z}^{(i)}$ and $\mathbf{z}^{(j)}$ should be close (**locality preserving**)
- $\mathbf{z}^{(i)}$ should allow for a good reconstruction of $\mathbf{x}^{(i)}$, i.e. there is some **decoder function** $g(\mathbf{z}^{(i)}) \approx \mathbf{x}^{(i)}$
- \mathcal{D}' captures as much **variance** or **information** as possible of \mathcal{D}

Applications of dimensionality reduction:

- lossy compression
- feature extraction (as input for other ML models)
- data visualization


Iris Setosa



petal

sepal

Iris Versicolor



petal

sepal

Iris Virginica



petal

sepal

Recall the **Iris dataset**, containing 4 features:

"sepalwidth", "sepalwidth", "petallength", "petalwidth", "class"

5.1, 3.5, 1.4, 0.2, Iris-setosa

4.9, 3.0, 1.4, 0.2, Iris-setosa

...

7.0, 3.2, 4.7, 1.4, Iris-versicolor



6.4, 3.2, 4.5, 1.5, Iris-versicolor

...


6.3, 3.3, 6.0, 2.5, Iris-virginica

5.8, 2.7, 5.1, 1.9, Iris-virginica

How to simultaneously plot 4 input features?



Computing the **first two principal components** transforms the data into a 2-dimensional space, which can be nicely visualized:



Principal Component Analysis: Formulation

Computing the PCA


Example: MNIST Digits

Linear Discriminant Analysis


Vectors $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_K$ are called **linearly independent**, when none of them can be expressed as a linear combination of the other vectors. That is, for each $1 \leq i \leq K$ and any coefficients z_k

$$\mathbf{x}_i \neq \sum_{k=1, k \neq i}^K z_k \mathbf{x}_k$$

Linearly dependent vectors




Linearly independent vectors



Let $B = (\mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_K)$ be a $D \times K$ -matrix with orthonormal vectors as columns, where $K < D$. The **linear subspace** spanned by B is defined as all possible linear combinations of B 's columns:

$$\mathbf{V} = \left\{ \mathbf{v} = B\mathbf{z} \mid \mathbf{z} = (z_1, \dots, z_K)^T \in \mathbb{R}^K \right\}$$



Let $B = (\mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_K)$ be a matrix with K orthonormal vectors as columns, spanning a K -dimensional subspace \mathbf{V} of \mathbb{R}^D .

We can **project** an arbitrary vector $\mathbf{x} \in \mathbb{R}^D$ onto \mathbf{V} by

- computing the **projection coefficients** $B^T \mathbf{x} =: \mathbf{z} \in \mathbb{R}^K$
- computing the **projection/reconstruction** $B\mathbf{z} =: \hat{\mathbf{x}} \in \mathbf{V}$
- thus, $\hat{\mathbf{x}} = \underbrace{BB^T}_{\text{projection matrix}} \mathbf{x}$
- $\hat{\mathbf{x}}$ is the **closest point** in \mathbf{V} (in Euclidean distance) to \mathbf{x}
- the **residual** $\mathbf{r} = \mathbf{x} - \hat{\mathbf{x}}$ is always orthogonal to $\hat{\mathbf{x}}$

Idea of PCA: Learn an “interesting” subspace \mathbf{V} .

Equivalently, learn a $D \times K$ -matrix B of orthonormal vectors capturing K “interesting directions.”


Principal Component Analysis: Formulation

Data

PCA is an **unsupervised** technique, i.e. we have a training dataset of feature vectors

$$\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

We can interpret the data as a “point cloud” in \mathbb{R}^D :



Step 1: Centering the Data


For PCA we assume that the average of the data is $\bar{\mathbf{x}} = 0$:

$$\bar{\mathbf{x}} := \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} = 0$$

If not, we simply remove the average from each training example:

$$\mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \bar{\mathbf{x}}$$


This effectively moves the center of the point cloud to the origin.




Thus, we will w.l.o.g. assume that $\bar{\mathbf{x}} = \mathbf{0}$.

Goal of PCA

- Find an “interesting subspace” \mathbf{V} of the data
- This amounts to finding a matrix $B = (\mathbf{b}_1, \dots, \mathbf{b}_K)$ of K orthonormal basis vectors
- The basis vectors \mathbf{b}_k define K “interesting directions”
- K is a **hyper-parameter** selected by the user
- For now, let's consider the special case $K = 1$, i.e. we want to find the “most interesting direction” in the data
- How to define “interestingness”?



Which of the three directions \mathbf{b} , \mathbf{b}' , \mathbf{b}'' is the most “interesting” one? Why?



\mathbf{b}' is the direction where

1. the **variance of the projected data is maximal**
2. the **sum of squared projection errors is minimal**

Let $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ be a centered training set of D -dimensional vectors.

Among all unit vectors, let \mathbf{b}_1 be the unit vector where

- the **variance of the projected data**
 $\text{var} \left(\left\{ z_1^{(1)}, z_1^{(2)}, \dots, z_1^{(N)} \right\} \right)$ is **maximal**, where

$$z_1^{(i)} = \mathbf{b}_1^T \mathbf{x}^{(i)}$$

or, **equivalently**, where

- the sum of **squared projection errors**

$$\sum_{i=1}^N \|\mathbf{x}^{(i)} - \hat{\mathbf{x}}^{(i)}\|_2^2 = \sum_{i=1}^N \|\mathbf{x}^{(i)} - \mathbf{b}_1 z_1^{(i)}\|_2^2$$

is **minimal**.

The vector \mathbf{b}_1 is called the **first principal direction** of dataset \mathcal{D} .

The values $\{z_1^{(1)}, z_1^{(2)}, \dots, z_1^{(N)}\}$ are called the **first principal component (scores)** of \mathcal{D} .

- the scores $z_1^{(1)}, z_1^{(2)}, \dots, z_1^{(N)}$ are a 1-dimensional representation of the original samples $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}$
- the vector $\hat{\mathbf{x}}^{(i)} = \mathbf{b}_1 z_1^{(i)} = \mathbf{b}_1 \mathbf{b}_1^T \mathbf{x}^{(i)}$ is the projection of $\mathbf{x}^{(i)}$ onto \mathbf{b}_1 (or, the best reconstruction of $\mathbf{x}^{(i)}$ given $z_1^{(i)}$)
- \mathbf{b}_1 is the **direction in which the data varies the most**, or the **direction of highest energy** – often, although not always, this direction carries the “most useful information” in the data

- the equivalence between **maximal variance** and **minimal projection/reconstruction error** (not proven here) is an example of **duality**, which occurs often in optimization
- by subtracting the projections $\hat{\mathbf{x}}^{(i)}$ from the original samples $\mathbf{x}^{(i)}$ and finding the principal component in the **residual vectors** $\mathbf{r}^{(i)} = \mathbf{x}^{(i)} - \hat{\mathbf{x}}^{(i)}$, one finds the **second principal component**, and so forth
- we will see in the following, however, that it is actually quite simple to find the first K principal components in one step

Computing the PCA

- Assume that the dataset $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ is already centered, i.e. $\bar{\mathbf{x}} = 0$
- We want to find the **first principal direction** \mathbf{b}_1
- To simplify notation, we denote it by \mathbf{b} , dropping the subscript

Problem Setup

- The PCA scores are given as

$$z^{(i)} = \mathbf{b}^T \mathbf{x}^{(i)}$$

- We wish to maximize the **empirical variance**

$$\text{var}(z^{(i)}) = \frac{1}{N} \sum_{i=1}^N (z^{(i)})^2 = \frac{1}{N} \sum_{i=1}^N (\mathbf{b}^T \mathbf{x}^{(i)}) (\mathbf{x}^{(i)T} \mathbf{b})$$

- We require that \mathbf{b} is a **unit vector**, i.e.

$$\mathbf{b}^T \mathbf{b} = 1, \quad \sum_{i=1}^D b_i^2 = 1$$

(note that if \mathbf{b} was not constrained to be a unit vector, $\text{var}(z^{(i)})$ would grow to infinity with the norm of \mathbf{b})

Problem Setup

We can transform the problem as follows:

$$\begin{aligned}\max_{\mathbf{b}^T \mathbf{b} = 1} \text{var}(z^{(i)}) &= \frac{1}{N} \sum_{i=1}^N z^{(i)} z^{(i)} \\ &= \frac{1}{N} \sum_{i=1}^N \mathbf{b}^T \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \mathbf{b} \\ &= \mathbf{b}^T \underbrace{\left(\frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)}_{\text{=: covariance matrix } C} \mathbf{b}\end{aligned}$$


Covariance Matrix

The (empirical) **covariance matrix** C is a $D \times D$ matrix containing all empirical (co-)variances in the data:

$$C = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} \mathbf{x}^{(i)T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}.$$

(**Recall** that \mathbf{X} is the design matrix, whose rows are the N training examples).

- the diagonal element C_{dd} is the **empirical variance** of the d^{th} dimension: $\frac{1}{N} \sum_i x_d^{(i)} x_d^{(i)}$
- the off-diagonal element C_{de} is the **empirical co-variance** between d^{th} and e^{th} dimension: $\frac{1}{N} \sum_i x_d^{(i)} x_e^{(i)}$
- C is always a symmetric matrix



PCA: Optimization Problem

With the covariance matrix, the PCA problem can compactly be written as

$$\max_{\mathbf{b}^T \mathbf{b} = 1} \mathbf{b}^T \mathbf{C} \mathbf{b}$$

What is the optimal \mathbf{b} here? Evidently, \mathbf{b} must be a vector somehow connected with \mathbf{C} .

- Key and solution to the PCA problem: **eigen decomposition** of C
- For any matrix C , when it holds for some vector \mathbf{v} and a scalar λ that

$$C\mathbf{v} = \lambda\mathbf{v}$$

then we call \mathbf{v} an **eigen vector** of C and λ its corresponding **eigen value**

- Without loss of generality, we assume that eigen vectors are normalized

For any **symmetric** $D \times D$ -matrix C , there are D **orthonormal** eigen vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_D$ and D sorted eigen values $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_D$, such that C can be written as

$$C = VEV^T$$

where $V = (\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_D)$ is an **orthonormal matrix** and E is diagonal:

$$E = \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_D \end{pmatrix}$$

For covariance matrices it further holds that all $\lambda_i \geq 0$; they are so-called **positive definite matrices**.

PCA Solution

Using the eigen decomposition we can write the PCA problem as

$$\max_{\mathbf{b}^T \mathbf{b} = 1} \mathbf{b}^T \mathbf{C} \mathbf{b} = \underbrace{\mathbf{b}^T \mathbf{V}}_{\mathbf{a}^T} \mathbf{E} \underbrace{\mathbf{V}^T \mathbf{b}}_{\mathbf{a}} = \sum_{i=1}^D \lambda_i a_i^2$$

- Since \mathbf{b} is a unit vector and \mathbf{V} is orthonormal, also \mathbf{a} is a unit vector, i.e. $\sum_{i=1}^D a_i^2 = 1$
- Evidently, $a_i^2 = (\mathbf{v}_i^T \mathbf{b})^2 \geq 0$
- Thus, the values a_i^2 are **non-negative** and **sum to one**
- The variance is

$$\lambda_1 a_1^2 + \lambda_2 a_2^2 + \cdots + \lambda_D a_D^2$$

- When is this maximal?

- The variance becomes maximal when $a_1^2 = 1$:

$$\lambda_1 1 + \lambda_2 0 + \cdots + \lambda_D 0$$


- Since $a_i = (\mathbf{v}_i^T \mathbf{b})$, this exactly happens if $\mathbf{b} = \mathbf{v}_1$!
- **Thus, the first eigen vector \mathbf{v}_1 of the covariance matrix is the first principal direction \mathbf{b}_1**
- The corresponding eigen value λ_1 is the variance which is “captured” in this direction (**explained variance**)

- More generally, the first K eigen vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_K$ are the the **first K principal directions**
- The sum of their corresponding eigen values is the variance explained by the first K principal components:

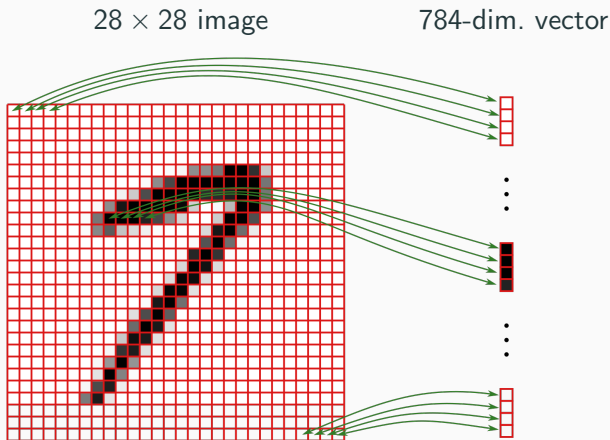
$$\text{"explained variance"} = \sum_{i=1}^K \lambda_i$$

Example: MNIST Digits

- collection of 60,000 images of handwritten digits
- images are 28×28 pixels, with 256 gray scales
- white = 0, black = 255




By following a fixed pixel order (e.g. row-wise) we can convert each image into a vector. By “uprolling” the vector, we get back an image.




The first step is to center the data, i.e. subtract the average vector (image) from each sample:

$$\bar{\mathbf{x}} = \frac{1}{60,000} \sum_{i=1}^{60,000} \mathbf{x}^{(i)}, \quad \forall i: \mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \bar{\mathbf{x}}$$

$\bar{\mathbf{x}} =$



$$C = \frac{1}{60,000} \sum_{i=1}^{60,000} \mathbf{x}^{(i)} \mathbf{x}^{(i)T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$



$$V, E = \text{eigen}(C)$$

The first K columns of V contain the first K **principal directions**.
We store them in a $784 \times K$ matrix B .

Note: Some implementations return the eigenvalues in arbitrary order.
In this case, sort them in descending order and sort the columns of V correspondingly.




black pixels correspond to positive values, white pixels correspond to negative values


The matrix B projects the data into a K -dimensional subspace:

$$\mathbf{Z} = \mathbf{X}\mathbf{B}$$


$K = 2$




$K = 3$



Reconstructions of four samples (first column) with 1, 10, 50, 100, 200 and 500 PCA components:



Reconstruction error (measured in mean squared error) and explained variance ($\sum_i \lambda_i$) over K :



→ can be used to select K

Given: dataset $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$, parameter K

1. Center data

- $\bar{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$
- for all i : $\mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \bar{\mathbf{x}}$

2. Compute covariance matrix

$$C = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)} \mathbf{x}^{(i)T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$

3. Compute eigen decomposition: $V, E \leftarrow \text{eigen}(C)$

4. Let B be the matrix consisting of the first K columns of V

5. **return** B

- The **principal component scores** (low-dimensional features) are then given as $\mathbf{Z} = \mathbf{X}B$ ($N \times K$ -matrix)
- Projections/reconstructions are given as $\hat{\mathbf{X}} = \mathbf{Z}B^T$
- Add $\bar{\mathbf{x}}$ again to the reconstructions


Linear Discriminant Analysis

Supervised Dimensionality Reduction

- PCA is **unsupervised**, i.e. there are no targets
- If one additionally has target values, one might use them to guide dimensionality reduction
- A classical method is Fisher's **linear discriminant analysis (LDA)**, using class information

Separating Classes

- Let $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$ be a binary classification dataset, i.e. each $y^{(i)} \in \{-1, 1\}$
- We want to learn a linear function $z = f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$ projecting D -dimensional vectors onto the real line
- This function should allow us to separate the two classes well
- How to select \mathbf{w} ?



Note: $\|\mathbf{w}\|$ is an arbitrary scaling, hence we can assume $\|\mathbf{w}\| = 1$.

Separating the Means?

First idea: Maximize the distance between projected means

- Let \mathbf{m}^- and \mathbf{m}^+ be the means of the negative and positive classes, respectively:

$$\mathbf{m}^- = \frac{1}{N^-} \sum_{i: y^{(i)} = -1} \mathbf{x}^{(i)} \quad \mathbf{m}^+ = \frac{1}{N^+} \sum_{i: y^{(i)} = +1} \mathbf{x}^{(i)}$$


Here N^- and N^+ are the number of positive and negative samples, respectively.

- Let m^- and m^+ be the projected means, i.e.

$$\mathbf{m}^- = \mathbf{w}^T \mathbf{m}^- \quad \mathbf{m}^+ = \mathbf{w}^T \mathbf{m}^+$$

- It can be shown, that $\mathbf{w} \propto \mathbf{m}^- - \mathbf{m}^+$, i.e. the vector pointing from one mean to the other, maximizes the distance between the projected means m^- and m^+

Separating the Means?



- The idea works to a certain extent, but apparently we are not doing a perfect job
- Considerable overlap between the classes in projected space
- The problem is that we ignored the **covariance structure** of the point clouds

- Let $z^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)}$ be the i^{th} projected sample
- Let again m^- and m^+ be the projected means
- Further, let v^- and v^+ be the **within-class variances** of the projected data:

$$v^- = \frac{1}{N^-} \sum_{i: y^{(i)} = -1} \left(z^{(i)} - m^- \right)^2 \quad v^+ = \frac{1}{N^+} \sum_{i: y^{(i)} = +1} \left(z^{(i)} - m^+ \right)^2$$

- The **Fisher criterion** is

$$J(\mathbf{w}) = \frac{(m^- - m^+)^2}{v^- + v^+}$$

- Increasing $J(\mathbf{w})$
 - increases the quadratic distance between m^- and m^+
 - and/or decreases variance (spread) within each class


- Fisher criterion $J(\mathbf{w})$ has an analytic maximum, denoted as **linear discriminant analysis (LDA)**
- Let C^- and C^+ be the covariance matrices of the negative and positive class in the original space

$$C^- = \frac{1}{N^-} \sum_{i: y^{(i)} = -1} (\mathbf{x}^{(i)} - \mathbf{m}^-)(\mathbf{x}^{(i)} - \mathbf{m}^-)^T$$
$$C^+ = \frac{1}{N^+} \sum_{i: y^{(i)} = +1} (\mathbf{x}^{(i)} - \mathbf{m}^+)(\mathbf{x}^{(i)} - \mathbf{m}^+)^T$$

- Then the maximum of $J(\mathbf{w})$ is

$$\mathbf{w}_{LDA} \propto (C^- + C^+)^{-1}(\mathbf{m}^- - \mathbf{m}^+)$$

In contrast to the previous solution, LDA perfectly separates the two classes in this example:



- Similar to PCA, LDA uses covariance structure to determine a linear subspace
- Main difference is that PCA is unsupervised, while LDA is supervised
- LDA can be generalized to many classes as well, yielding a $(|\mathcal{C}| - 1)$ -dimensional subspace